

Neural Progressive Photon Mapping

paper1053

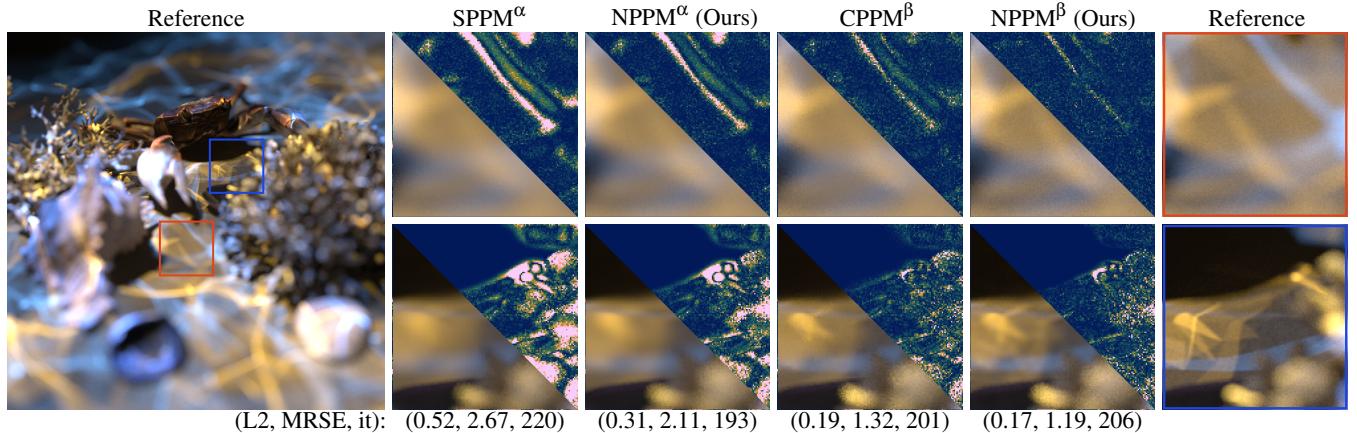


Figure 1: Equal-time comparison (15 seconds) between our neural progressive photon mapping (NPPM), SPPM [HJ09] and CPPM [LLZ^{*}20] on the CRAB DOF scene with depth of field. The superscripts α and β respectively refer to the different radius reduction policy used by the two baseline methods, which we incorporate atop NPPM. Our technique reduces the overall bias compared to its nonneural counterparts, capturing sharper caustics on most of the scene. False color error shows the MRSE metric.

Abstract

Photon density estimation is a robust solution for estimating complex light transport, such as those involving caustics and pure specular interactions. The shape and bandwidth of the density kernel are both crucial in achieving optimal performance. Recently, density kernels directly predicted by neural networks from local photon statistics have shown improved reconstruction results for small numbers of photons. The direct weight prediction approach of these methods, however, is fundamentally incompatible with consistent estimators as it does not allow for direct control over bias and variance. We address this problem by relying on a simpler yet effective analytical kernel, also inferred by a neural network. Unlike prior work, our technique supports progressive schemes by design, hence unlocking a large variety of applications such as stochastic photon mapping. Our method is fast, trivial to train and demonstrates state-of-the-art caustics reconstruction at equal-time over other photon mapping techniques.

CCS Concepts

- Computing methodologies → Ray tracing; Neural networks;

1. Introduction

Light transport simulation aims to recreate the complex interactions between light and matter in virtual environments. These interactions, especially *caustics*—like light focusing through a whisky glass or sunrays slowly dancing on coral reefs—are crucial to photorealism, especially in architectural visualization and visual effects. Caustic light paths formed from long specular chains present unique challenges to classical Monte Carlo algorithms: not only are they difficult to discover but they are also hard to explore efficiently once found. Consequently, specialized rendering techniques

are required to achieve acceptable noise levels in production. Such methods focus, for instance, on modeling and solving the intricate geometric constraints of such paths [ZGJ20, FGW^{*}24, FWW^{*}25].

An alternative approach to traditional methods like path tracing is to employ spatial relaxation [Jen96], where the physical constraints of specular chains are softened to enable path reuse through nearby clustering. Density estimation represents one such family of techniques and offers an attractive balance between implementation simplicity and robustness. A significant practical challenge, however, lies in determining the appropriate spatial relaxation kernel

parameters. This in turn introduces an inherent trade-off between kernel shape and bandwidth, where one must maximize path reuse while also minimizing the bias introduced in the underlying estimator. As a result, the process of selecting optimal kernel settings tends to be laborious and context-dependent as it must account for local and arbitrarily complex radiometric properties such as photon distribution patterns in a given scene.

To mitigate this problem, prior works have proposed adaptive and progressive density estimation methods based on statistics gathered during the rendering pass [KD13, SFES07, KWX*16, PAMM19]. In particular, deep photon mapping (DPM) [ZXJ*20] employs K -nearest neighbor (K -NN) and a small neural network to predict weights for each photon, demonstrating impressive results over standard photon mapping approaches. Intuitively, DPM learns to adapt to the local context by exploiting statistics from collected nearby photons. Unfortunately, this method cannot be combined with consistent estimators as the kernel bandwidth and bias cannot be directly controlled. This lack of control prevents applications to a wider range of rendering methods [MBGJ22, WGH22, HPJ12, GKDS12].

In light of these observations, we introduce neural progressive photon mapping (NPPM), a learning-based framework that predicts spatially adaptive kernels with progressive rendering in mind. Our network generates *parametric and normalized* kernels that can be restricted at will during the iterative rendering process. Since we do not predict per-photon weights, our architecture eliminates the need to compute and store deep context features for each photon. Consequently, our method is not only simpler than DPM but also fully controllable. Moreover, our design allows the use of a 3D grid to average latent photon representations over time, enabling further optimizations. Finally, we show how to effectively train our model while ensuring compatibility with radius range queries, hence introducing the first learning-based *progressive* photon mapping method to the rendering community.

To summarize, our contributions are:

- a simple and lightweight parametric neural model that can efficiently adapt to the local photon distribution;
- a controllable density kernel scheme compatible with progressive rendering methods; and
- an expressive latent representation that can be averaged and stored on the scene surfaces.

2. Related Work

Photon mapping. Jensen's seminal work on density estimation [Jen96] has shown great promise in rendering complex light phenomena like caustics. Photon mapping (PM) is now standard in light transport simulation and has been adopted by production renderers like the Corona renderer [ŠK19]. PM has been successfully adapted to BSSRDF rendering [JMLH01], volume photon mapping [JC98], and beam radiance estimation [JZJ08], and can also be combined with other path construction techniques [GKDS12, HPJ12, KGH*14] via multiple importance sampling [Vea98] for improved robustness.

A key consideration in practice is the bias-variance trade-off

of photon mappers. The amount of bias can be reduced by incorporating ray differential information [SFES07], local radiance gradient estimates [KWX*16], more flexible (but generally more costly) kernels [PAMM19] or progressive kernel bandwidth reduction schemes [HOJ08, HJ09, KZ11]. Since these kernel reduction techniques can be sensitive to a nonoptimal initial radii, the overall convergence of the estimator may suffer dramatically, leading to either excessive bias or variance in the final image. A mitigation measure is to derive improved progressively shrinking strategies based on derivative information [KD13] or χ^2 -statistics [LLZ*20]. We focus on enabling progressive estimation through large, controllable kernels while simultaneously minimizing bias. Our method is compatible with different radius reduction policies and surpasses the performance of prior works.

Note that bias in density estimation can be further removed based on Bernoulli trials [QSH*15] or telescopic series [MBGJ22], at the cost of a slight increase in variance. We see these techniques as orthogonal to ours and we expect our method to be compatible due to our controllable kernel support. This control potentially allows our method to be applied to path filtering algorithms, which perform local gathering with range queries [KDB16, WGGH20] or local clustering [DHC*21], to reduce variance. Alternatively, methods for guiding photon's emission [GPGSK18, EK20] or their entire path [HJ11, CWY11, GRŠ*16, vOHK16] can similarly be used in our method to improve the overall photon distribution.

Neural methods for density estimation. Learning-based approaches have proven successful in many areas of rendering, including Monte Carlo image denoising [CKS*17, IMF*21, HY21]. Such reconstruction techniques can be adapted to density estimation [CM21] by extracting and combining correlated images with different biases. We can incorporate these techniques into NPPM, as we can bound the bias we introduce.

Photon mapping has also been applied to path guiding by binning photon contributions [Jen95]. Recently, neural methods have been used to denoise these distributions [ZXS*21b], even when they are stored in a quadtree [ZXS*21a]. Such methods are generally more memory intensive and require complex spatial partitioning strategies. Alternatively, neural networks can directly predict distribution in path guiding [MMR*19] based on the framework of normalizing flows [KPB21]. However, due to their complex architecture, these approaches are often too expensive except when targeting lower-dimensional distributions [LHL*24]. Simpler, online neural guiding methods can also optimize for mixtures of von Mises–Fischer [DWL23] or anisotropic Gaussian [HIT*24] distributions. Our work is inspired by these parametric models as they are both more computationally efficient and expressive enough for density estimation.

To the best of our knowledge, only Zhu et al. [ZXJ*20] have employed neural networks for density estimation in rendering. Their proposed DPM approach leverages a PointNet-like architecture [QSMG17] to directly regress photon weights, resulting in neural kernels that outperform previous non-neural density estimators at different photon sampling rates. DPM does not support a varying number of nearest neighbors and, importantly, cannot be directly used in a progressive estimator. NPPM addresses both these problems by providing an estimator that is consistent by construction.

130 **3. Background**

131 **Density estimation.** In surface rendering, we are interested in
132 computing the exitant radiance L_o at a location \mathbf{x} in the outgoing
133 direction ω_o :

$$134 \quad L_o(\mathbf{x}, \omega_o) = \int_{\Omega} f_r(\mathbf{x}, \omega_o, \omega) L_i(\mathbf{x}, \omega) (\mathbf{n} \cdot \omega)^+ d\omega, \quad (1)$$

134 where f_r is the BRDF, L_i the incoming radiance in direction $\omega \in \Omega$,
135 and \mathbf{n} is the surface normal at \mathbf{x} . As proposed in classical photon
136 mapping [Jen96], this exitant radiance can be formulated as a ker-
137 nel density estimation (KDE) problem on the photon map:

$$138 \quad \langle L_o \rangle_N^{\text{PM}} = \frac{1}{N} \sum_{n=1}^N Q_r(\|\mathbf{x}_n - \mathbf{x}\|) f_r(\mathbf{x}, \omega_o, \omega_n) \Phi_k, \quad (2)$$

138 where N is the total number of emitted photons, r is the search ra-
139 dius, \mathbf{x}_n is the position of the n -th photon with flux Φ_n and incoming
140 direction ω_n . The flux includes the throughput of the photon. The
141 normalized kernel Q has support within radius $r > 0$ only, allowing
142 fast spatial query through specialized data structure or by using ray
143 tracing hardware routines [KBG23]. In general, this spatial kernel
144 is chosen to be simple and radial; for example $Q_r(d) = (\pi r^2)^{-1}$ if
145 $d < r$ otherwise 0.

146 **Deep density estimation.** Zhu et al. [ZXJ*20] replace paramet-
147 ric kernels with a neural network that directly regresses a weight
148 for each collected photon. Under this framework, a deep context
149 vector $\mathbf{c} \in \mathbb{R}^c$ is inferred from the nearby collected photons using
150 K -nearest neighbors at the gather point location \mathbf{x} :

$$151 \quad \rho_k(\mathbf{x}) := [\mathbf{x}'_k, \omega_k, \Phi'_k], \\ 152 \quad \mathbf{c}(\mathbf{x}; \Theta) = \text{Pool}_K \left(\{F(\rho_k(\mathbf{x}); \Theta)\}_{k=1}^K \right). \quad (3)$$

151 Here, ρ_k represents the concatenated photons' input features in-
152 cluding the normalized position within the kernel support trans-
153 formed in tangent space \mathbf{x}'_k , the photon's incoming directions ω_k
154 and the associated tonemapped flux Φ'_k . Each photon is encoded
155 independently by a feature extractor network F with trainable pa-
156 rameters Θ . To ensure input order invariance, the deep context vec-
157 tor is computed by average/max pooling across all K photons, fol-
158 lowing the design principles of PointNet [QSMG17]. The resulting
159 deep context vector (DCV) is then decoded by a separate network
160 D with trainable parameters Φ , substituting the kernel function in
161 Eq. (2) to yield the estimator:

$$162 \quad \langle L_o \rangle_K^{\text{DPM}} = \frac{1}{N\pi r^2} \sum_{k=1}^K w_k f_r(\mathbf{x}, \omega_o, \omega_k) \Phi_k, \quad (4)$$

$$163 \quad w_k = D([\rho_k(\mathbf{x}), \mathbf{c}(\mathbf{x}, \Theta)]; \Phi).$$

162 A per-photon weighting scalar $w_k > 0$ is predicted and the nor-
163 malization factor $(\pi r^2)^{-1}$ ensures that it remains scale-invariant with
164 respect to the K -NN radius.

165 Both networks are small multilayer perceptrons (MLPs) that are
166 trained jointly to minimize the L^2 -norm between the estimated ra-
167 diance from Eq. (4) and the target radiance $L_o(\mathbf{x}, \omega_o)$ obtained from
168 converged references. The error can stem from two sources: a poor
169 kernel (e.g., geometric discontinuities within its support), which

170 DPM is designed to handle, or too few photons to suppress vari-
171 ance. In the latter case, the incompatibility of DPM with the pro-
172 gressive frameworks blocks further variance reduction.

173 **Progressive density estimation.** Progressive photon mapping
174 (PPM) [HOJ08, HJ09, KZ11] introduces the concept of averaging
175 density estimates through multiple iterations. At each iteration,
176 the radius r is reduced to form a consistent estimator. Knaus and
177 Zwicker [KZ11] propose with APA to average the different density
178 estimates at each pixel and update the kernel's radii as

$$179 \quad r_{i+1}^2 = \frac{i + \alpha}{i + 1} r_i^2, \quad i \in \mathbb{Z}_{\geq 0}, \quad (5)$$

180 where $\alpha := 2/3$ is a user-defined parameter to balance bias reduc-
181 tion and variance increase. The initial radius r_0 can be set to a con-
182 stant user input, optionally scaled by the scene's extents, camera
183 ray differential or K -NN estimates on an independent photon map.
184 Alternate radius reduction schemes also exist [KD13, LLZ*20]. We
185 also utilize CPPM's radius reduction policy which applies a thresh-
186 old on the minimum number of gathered photons ($N_{\min} = 10$) be-
187 fore reducing the radius by a constant factor ($k = \sqrt{0.8}$). To ensure
188 a consistent estimator, the minimum photon count is then increased
189 at each radius reduction step by a constant factor $\beta = 1.2$. We add
190 subscripts NPPM $^\alpha$ and NPPM $^\beta$ to denote which radius reduction
191 strategy (APA's or CPPM's, respectively) is used in our results.

192 **3.1. Motivation**

193 The application of photon density estimation often involves a deli-
194 cate balance between minimizing both variance and bias. Reducing
195 variance can be achieved by averaging more photons per gather
196 point by increasing the photon density via more aggressive photon
197 tracing or by using larger density kernels. These solutions can be
198 impractical and inefficient when implemented without extra care:

- **Larger kernels** effectively increase the number of reused paths
199 at each gathering point, amortizing the cost of the photon trac-
200 ing pass. However, these kernels can introduce additional bias if
201 their shapes do not adapt to the underlying photons' density pro-
202 files. DPM address this by directly learning to predict the pho-
203 ton's weights. This relies on K -NN for collecting photons and
204 does not offer a mechanism to compute the kernel normaliza-
205 tion factor. This is problematic when trying to control the bias
206 introduced.
- **Tracing more photons** naturally leads to memory problems. To
207 solve this, iterative methods can be employed to average pho-
208 ton passes. This iterative strategy can be combined with radius
209 reduction to make the estimator consistent.

210 Our main objective is thus to design a deep adaptive kernel that
211 can be combined with iterative radius reduction. To this end, we
212 need our adaptive kernel to be *controllable*, implying we can ex-
213 press it at *any* bandwidth value. This in turn requires our method to
214 support range queries rather than K -NN. Our method must also be
215 lightweight as a massive amount of kernel density estimates will be
216 performed throughout the rendering process. Finally, our method
217 must retain information across iterations to ensure stable and pre-
218 cise predictions over time.

220 4. Neural progressive photon mapping

221 We now introduce our neural progressive photon mapping (NPPM)
 222 algorithm and detail how consistency is enforced. Similar to DPM
 223 [ZXJ^{*}20], our approach modifies only the kernel evaluation, leaving
 224 the rest of the photon mapping architecture unchanged. We first
 225 describe the model architecture, including how our adaptive kernel
 226 is normalized (Section 4.1). Next, we present a practical applica-
 227 tion of our framework using 2D anisotropic Gaussian kernels (Sec-
 228 tion 4.2). Finally, we discuss practical considerations (Section 4.3).

229 4.1. Model architecture

230 Figure 2 illustrates our framework. Similarly to DPM, we first pre-
 231 process photon features to construct the input to our model. We
 232 choose $\rho_k(\mathbf{x}) = [\mathbf{x}'_k]$ only. We express \mathbf{x}'_k in tangent space (with
 233 $z = 0$ for the surface plane) and use only this feature, as adding
 234 photon directions or flux did not improve convergence in our ex-
 235 periments. For instance, applying Russian roulette, photon flux re-
 236 mains nearly constant across photons, rendering this feature unin-
 237 formative. We also discard the z component of \mathbf{x}'_k , since retaining it
 238 provided no measurable benefit. We encode each photon’s statistics
 239 using a lightweight feature extractor network F . These latent fea-
 240 tures are subsequently aggregated through both average/max pool-
 241 ing layers to construct the deep context vector \mathbf{c} , effectively rep-
 242 resenting the gathered photon statistics in a shared latent space.
 243 We augment this representation with a global statistic v (described
 244 in Section 4.3), which encodes the overall photon density within
 245 the kernel. A second network D then predicts our parametric ker-
 246 nel. While this applies to arbitrary kernel families, we show that
 247 iso/anisotropic Gaussian densities are sufficient candidates.

248 **Kernel normalization estimator.** Based on the predicted kernel
 249 parameters, we compute a normalization factor to control the bias
 250 in our method. We restrict the kernel domain to maintain this con-
 251 trol by using the 2D disk \mathcal{B}_r of radius $r > 0$. However, using a given
 252 world-space radius makes our method scale-sensitive and hinders
 253 its generalization across different scenes. We address this by refor-
 254 mulating our adaptive kernel Q in a normalized space where $r = 1$.
 255 This reparametrization yields the following normalization factor

$$290 Z := \left(\int_{\mathcal{B}} Q(\mathbf{x}; \mathbf{c}, \boldsymbol{\varphi}) d\mathbf{x} \right)^{-1}, \quad (6)$$

256 which reciprocal can be estimated via Monte Carlo integration as

$$301 \langle Z^{-1} \rangle_M = \frac{1}{M} \sum_{m=1}^M \frac{Q(\mathbf{x}_m; \mathbf{c}, \boldsymbol{\varphi})}{p(\mathbf{x}_m)}, \quad \mathbf{x}_m \sim p, \quad (7)$$

257 where p denotes the probability density function. Note that, be-
 258 cause we estimate the reciprocal, Eq. (7) can introduce bias in the
 259 true normalization factor’s estimate. This bias can be reduced to an
 260 negligible level with sufficient samples or by applying debiasing
 261 techniques [QSH^{*}15, MBGJ22]. Since both methods incur signifi-
 262 cant overhead, a more practical solution is to precompute and tab-
 263 ulate the normalization factor when it is smooth and the kernel has
 264 few parameters.

265 **Parametric density estimator.** We use our normalized adaptive
 266 kernel to perform density estimation. Writing $\mathbf{d}_k := \mathbf{x}_k - \mathbf{x}$ as the

267 2D displacement with respect to the gather point, our estimator is

$$\langle L_o \rangle_K^{\text{NPPM}} = \frac{Z}{N\pi r^2} \sum_{k=1}^K Q(\mathbf{d}_k/r) f_r(\mathbf{x}, \omega_0, \omega_k) \Phi_k, \quad (8)$$

268 where the factor πr^2 accounts for the radius rescaling. Note that
 269 our estimator is similar to the KDE formulation on the photon
 270 map (Eq. (2)). More importantly, it differs from DPM’s estimator
 271 (Eq. (4)) in two key aspects:

- 272 • We decode the parametric kernel values using the DCV *only once*. This approach is significantly more efficient than performing
 273 an MLP inference for each photon. This also allows our
 274 method to amortize the cost of computing the DCV by exploiting
 275 its spatial redundancy in 3D.
- 276 • Our kernel’s bias is controllable by restricting its domain. In ad-
 277 dition, the kernel can be easily normalized through precompu-
 278 tation. This eliminates the need to jointly learn the kernel’s shape
 279 and its normalization, enabling faster and more stable training.

281 4.2. Application: 2D truncated anisotropic Gaussian kernel

282 **Kernel formulation.** A suitable kernel candidate for our method
 283 is a 2D anisotropic Gaussian centered at the gather point. In the
 284 normalized space, this Gaussian has zero mean. Inspired by the
 285 formulation from 3D Gaussian Splatting [KKLD23], we form a co-
 286 variance matrix Σ by constructing a 2D scaling matrix $S(\mathbf{c}, \boldsymbol{\varphi})$ and
 287 rotation matrix $R(\mathbf{c}, \boldsymbol{\varphi})$:

$$\Sigma(\mathbf{c}, \boldsymbol{\varphi}) = RSS^T R^T, \quad (9)$$

288 ensuring that Σ is symmetric and positive semi-definite. Our pre-
 289 dicted kernel is then defined as

$$Q_{\Sigma}^{\text{2DG}}(\mathbf{d}) = \mathbf{1}_{\mathcal{B}}(\mathbf{d}) \cdot \mathcal{N}(\mathbf{d}; \mathbf{0}, \Sigma) / Z, \quad (10)$$

290 where $\mathbf{1}_{\mathcal{B}}$ is the indicator function on the unit disk and we have
 291 dropped the dependency on the deep context vector and the de-
 292 coder parameters for simplicity. Intuitively, our kernel resembles
 293 an anisotropic Gaussian that gets truncated at unit radius (Fig. 2).

294 **Kernel normalization.** Since our predicted covariance matrix Σ
 295 is determined by three parameters but the reciprocal normaliza-
 296 tion factor (Eq. (7)) is invariant under rotation due to the disk-
 297 shaped domain \mathcal{B} , the look-up table (LUT) only needs to be two-
 298 dimensional. We reparametrize our scale in log-space to get more
 299 precision for small scale values. The full LUT takes at most 4 MB
 300 of GPU memory.

301 **Kernel hyperparameters.** We construct our rotation matrix R by
 302 predicting an angle ξ and rescaling by 2π . The scale matrix S is built
 303 by predicting two scaling factors $s_1, s_2 > 0$. To ensure numerical
 304 stability and prevent degenerate scales, we use an ELU activation
 305 with $\alpha = 1$ and $\epsilon = 1.0001$:

$$306 s_i = \epsilon + \text{ELU}(z_i - 2) = \epsilon + \begin{cases} z_i - 2, & \text{if } z_i \geq 2 \\ \exp(z_i - 2), & \text{if } z_i < 2 \end{cases} \quad (11)$$

307 with z_i being the i -th output of the network. We clamp the scale
 308 value to a maximum value of $\sqrt{3}$, which makes the predicted Gaus-
 309 sian corresponds roughly to a box kernel. We fix these hyperpara-
 310 meters across all experiments.

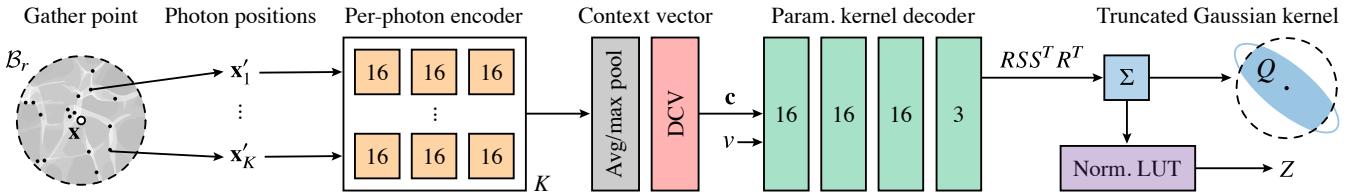


Figure 2: Our NPPM architecture. We gather nearby photons in a given radius and compute per-photon features using a lightweight MLP encoder F . The resulting features are then pooled into a deep context vector (DCV). This embedding is consumed by a decoder network D , alongside the gathered density statistics, to produce the parameters (rotation and scale) of the 2D covariance matrix of a truncated Gaussian distribution. Thanks to the symmetric nature of this kernel, we can precompute the corresponding normalization factor up to a rotation in a small look-up table, hence enabling exact density estimation.

310 4.3. Practical considerations

311 Implementing our algorithm efficiently requires addressing several
 312 key considerations. This section examines how to support range
 313 queries during both training and inference phases of our method.
 314 We also demonstrate how our network design facilitates caching
 315 and reuse of latent representations across different rendering it-
 316 erations and scene locations, hence enhancing computational ef-
 317 ficiency.

318 **Range versus nearest neighbor queries.** DPM [ZXJ*20] uses K -
 319 NN since it does not need to control bias. Additionally, K -NN sim-
 320 plifies implementation by ensuring that a constant number of pho-
 321 tons is processed by the network, making both training and infer-
 322 ence more parallelizable: since the number of photons per gather
 323 point is fixed, the memory can be preallocated efficiently.

324 Our method is also compatible with nearest neighbors and can be
 325 trained using this density estimation strategy. As our domain is re-
 326 stricted to a disk \mathcal{B}_r however, we require range queries instead. This
 327 results in a *variable* number of input photons, which complicates
 328 batch training. To address this, we impose a fixed upper bound on
 329 the number of photons collected. This leads to two scenarios:

330 **I. Fewer photons than the maximum:** We mask the missing pho-
 331 tons by zeroing out their encodings and setting their associated
 332 weights in the predicted kernel to zero. This introduces bias in
 333 the average pooling. To compensate, we apply a correction fac-
 334 tor $\zeta = B/K_{\text{matches}}$ to the mean vector, where B is the fixed batch
 335 size (i.e., the maximum number of photons) and K_{matches} is the
 336 actual number of photons gathered. Max-pooling remains unaf-
 337 fected by this change.

338 **II. More photons than the maximum:** We uniformly resample the
 339 photons to match the desired count. This can be efficiently im-
 340 plemented using reservoir sampling, avoiding the need to store
 341 the full list of gathered photons. We also store ζ as it is an eas-
 342 ily accessible and valuable statistic to inform our decoder of the
 343 relative photon density at the gather point location.

344 **Computing the deep context vector** In addition to the generated
 345 deep context vector \mathbf{c} , we append a compressed representation of
 346 the relative photon density, defined as $v = \log(1 + \zeta)$, to avoid sat-
 347 urating \mathbf{c} . This statistic is inexpensive to obtain, varies smoothly
 348 across the scene and can guide the decoder towards more aligned
 349 kernels.

350 **Deep kernel evaluation** After constructing the deep context vec-
 351 tor \mathbf{c} with v , we decode the parameters of our normalized para-
 352 metric kernel. Using these parameters, we query the correspond-
 353 ing normalization factor from our LUT and evaluate the KDE esti-
 354 mator (Eq. (8)). This estimator can either be applied to the (resam-
 355 pled) subset of photons used to form the network input or to all
 356 photons within the gather point radius via a separate range query.
 357 Both approaches have trade-offs: reusing the resampled photons
 358 reduces computation, but increases variance, especially when the
 359 predicted kernel is narrow. We found that performing an additional
 360 range query is more advantageous in practice, as it does not intro-
 361 duce extra variance and its cost can be amortized using a spatial
 362 latent averaging strategy, which we describe next.

363 4.4. Spatial latent averaging

364 Our approach naturally allows for decoupling the encoding and de-
 365 coding of the deep context vector. To this end, we maintain a la-
 366 tent representation at each pixel and average these representations
 367 over multiple iterations. This approach allows the kernel’s adaptive
 368 properties to gradually decay as rendering proceeds. While the en-
 369 coder network is most computationally expensive (as it is applied
 370 independently to each of the B photons selected per gather point),
 371 the decoding network is invoked only once per gather point. As
 372 such, an opportunity arises to amortize the DCV encoding opera-
 373 tion across gather points and iterations.

374 By exploiting the smoothness of the photon distribution over the
 375 scene surfaces, we propose to combine our method with a 3D grid
 376 that store and access these latent representations. **The grid size is
 377 initialized according to the largest axis of the scene’s bounding box.**
 378 At each iteration, we perform the following two steps:

- 379 1. Collect photons at a subset of gather points, encode them, and
 380 average their resulting DCVs along with the density statistic v
 381 inside the 3D grid. To avoid aliasing, we apply stochastic jitter-
 382 ing to the gather point locations before splatting [BFK20].
- 383 2. Retrieve the averaged deep context vector from the grid, decode
 384 the corresponding adaptive parametric kernel, and perform ker-
 385 nel density estimation using a range query (Eq. (8)).

386 Averaging the spatial DCV can be done using a small exponential
 387 smoothing factor ($\gamma = 0.5$) or by averaging over all iterations. We
 388 found that averaging over all iterations gives slightly better results
 389 than exponential smoothing. We leave the investigation on the best



Figure 3: Examples of our procedurally generated training dataset. Gather points travelling through dielectric surfaces are masked out to avoid excessive noise during training.

390 policy to update our spatial grid as future work. This grid approach
 391 achieves better performance compared to predicting the DCV at
 392 every iteration for all gather points, as it reduces noise in both the
 393 produced DCV and the density statistic v , as well reduce our overall
 394 cost. Sharp discontinuities in the photon density can introduce bias;
 395 however, this bias remains bounded due to the underlying properties
 396 of our kernel and will vanish during the progressive rendering
 397 process.

398 **Early stopping** To further reduce overhead, we stop updating the
 399 latent representations after 50 iterations. This strategy is similar to
 400 that used in neural path guiding [DWL23, HIT^{*}24]; however, unlike
 401 neural path guiding—which relies on online training and maintaining
 402 computation graphs—our method only performs MLP evaluations
 403 per photon for encoding.

404 **Surface normal test** During photon gathering (both for producing
 405 the DCV or performing KDE (Eq. (8)), we reject photons whose
 406 surface normals significantly deviate from that of the gather point.
 407 This filtering strategy is common in traditional photon mapping to
 408 reduce bias. In our setting, we found that applying this normal-
 409 based filtering helps the network focus on more relevant regions of
 410 the scene, rather than over-emphasizing geometric discontinuities
 411 such as edges, creases and corners. Our adaptive kernels account
 412 for these regions by predicting better suited kernels.

413 4.5. Training

414 **Dataset generation.** Our training set consists of 50 procedurally
 415 generated scenes rendered at a resolution of 256×256 . Figure 3
 416 shows a few representative examples. In all scenes, we isolate and
 417 compute only the caustic component of light transport to let the
 418 network focus on complex and high-frequency photon distributions
 419 typical of caustics. Reference images are generated using progres-
 420 sive photon mapping [HOJ08]. Range queries are performed at run-
 421 time to augment our data with random photon maps. We uniformly
 422 pick radii in fixed intervals to add variability in the density of pho-
 423 tons. We set the maximum number of photon to $K_{\max} = 1000$ and
 424 apply our masking and resampling strategy from Section 4.3. We
 425 use a high value here to mitigate the noise on training gradients
 426 caused by the resampling step. In total, our dataset is composed of
 427 more than 3.2 million supervised training samples.

428 The scenes themselves are randomly generated from an interior
 429 cube. We assign random albedo colors to each diffuse wall and
 430 place a rectangular area light source with random orientation, scale,
 431

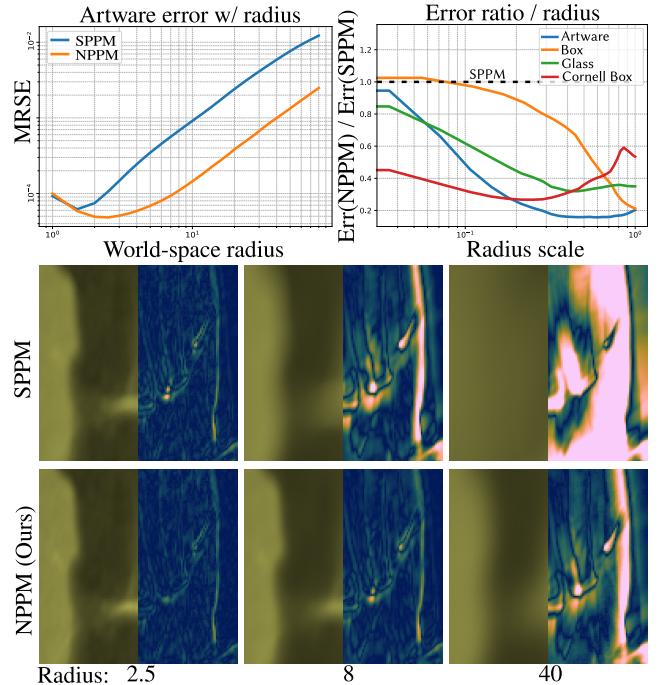


Figure 4: We plot and visualize the MRSE error for SPPM and our method across various radius scales. Our method consistently outperforms the box kernel, with diminishing returns at smaller radii.

432 and position. We then insert multiple dielectric for caustics and dif-
 433 fuse objects for shadows and non-flat surfaces shading scenarios.
 434 We apply random bump mapping to all dielectrics, producing a
 435 wide variety of intricate light patterns. We further randomize the
 436 position and look-at direction of the camera.

437 **Encoder pre-training.** To stabilize and accelerate training, we
 438 pretrain our encoder network $F(\theta)$ within a variational autoencoder
 439 (VAE) [KW14]. The VAE is trained to capture and reconstruct the
 440 input photon distribution, thereby promoting an expressive deep
 441 context vector (DCV) that generalizes across a wide range of pho-
 442 ton distributions. To reuse the pretrained encoder in our kernel
 443 adaptation network (Fig. 2), we predict the mean and variance of
 444 $K = 32$ Gaussians through a single linear projection from the DCV.
 445 New photons are then generated by sampling from these Gaussians
 446 and decoding their positions with a small MLP. The VAE is opti-
 447 mized for 300 training steps on small subsets of the dataset by
 448 minimizing the Chamfer distance between the generated and input
 449 photons. In practice, we found it more effective to freeze the en-
 450 coder and use it directly to produce the DCV fed into our kernel
 451 decoder $D(\phi)$, rather than retraining it from scratch.

452 **Model hyperparameters and optimization.** Our feature extrac-
 453 tor network $F(\theta)$ is a 3-layer MLP with Tanh activations, while
 454 our decoder $D(\phi)$ shares the same characteristics. We use $c = 32$
 455 dimensions for our deep context vectors.

456 We use the Adam optimizer [KB15] with a learning rate of 10^{-4} .
 457 To manage memory usage during training, we subsample 50% of
 458

457 the image pixels, as backpropagation over all gathered photons ex-
 458 ceeds our GPU memory capacity. We minimize the L^2 -loss com-
 459 puted between the predicted and target radiance values, both passed
 460 through the compression function

$$L \mapsto \log(1 + \mu L) / \log(1 + \mu), \quad (12)$$

461 where $\mu = 5000$, following Zhu et al. [ZXJ*20]. We initialize the
 462 network to produce small kernels at the start of training, as we
 463 observed that large initial kernels often cause the optimization to
 464 converge to poor local minima. We optimize our model for only
 465 10 epochs on an NVIDIA RTX 3080 GPU. This takes roughly 90
 466 minutes, a stark contrast with DPM reporting two days of training.

467 5. Results

468 **Implementation details.** We use Mitsuba 3 [JSR*22] to perform
 469 the eye pass and the photon tracing pass. We then execute our radius
 470 search on the GPU using a custom CUDA implementation of fixed-
 471 radius nearest neighbors search by Rama [Ram14]. We use Py-
 472 Torch [PGM*19] and fully-fused MLPs [M21] for acceleration during
 473 training. All our test references are computed with stochastic
 474 progressive photon mapping (SPPM) [HJ09] to have fresh gather
 475 points at every iterations. We reimplemented DPM [ZXJ*20] as
 476 no open source codebase is available. Similarly, we reimplemented
 477 CPPM [LLZ*20] in Mitsuba 3. All learning methods are trained on
 478 our dataset described in Section 4.5. We will release our implemen-
 479 tation and our generated dataset upon publication.

480 To evaluate our approach we use Mean Squared Error (L^2) and
 481 Mean Relative Squared Error (MRSE) with $\epsilon = 10^{-2}$. We also used
 482 the Symmetric Mean Absolute Percentage Error (SMAPE) to eval-
 483 uate fireflies in the unbiased application. All results were obtained
 484 under equal iteration counts or equal runtime conditions. For the
 485 photon encoding, we use range query with 50 maximum photons.
 486 SPPM, CPPM, and our NPPM use all photons. We trace 2M pho-
 487 tons/iteration for the CRAB DOF and CRAB scenes (Figs. 1, 6
 488 and 11); all other scenes use 400K photons/iteration. For glossy
 489 materials, we use 0.1 as minimum roughness to deposit photons and
 490 gather points. No guiding is performed during the tracing. DPM
 491 contains a total of 6.2K parameters, whereas ours only has 2.5K
 492 parameters.

493 **Sensitivity to range query size.** Figure 4 shows how our method
 494 is more robust to the scale of the range query compared to the non-
 495 adaptive box kernel from SPPM. All closeup images are produced
 496 by averaging multiple density estimates to make the bias more ap-
 497 parent. The improvement of our method diminishes with smaller
 498 kernels, unsurprisingly, as the assumption of uniform density holds
 499 in this regime. When the radius falls below a certain threshold (1.0
 500 ray differential radius), we switch to a box kernel to maintain nu-
 501 matical stability and reduce computational cost.

502 **Consistent method comparisons.** Figure 6 shows the comparison
 503 between our methods with anisotropic 2D kernels, SPPM [HJ09]
 504 and CPPM [LLZ*20]. We omit DPM [ZXJ*20] since it cannot be
 505 made consistent. CPPM has a different reduction policy based on
 506 the χ^2 -test, which often reduces the radius faster than SPPM, hence
 507 reducing the overall bias. Our method NPPM $^\alpha$ uses the same re-
 508 duction policy as APA, so that each iteration gather the same set of

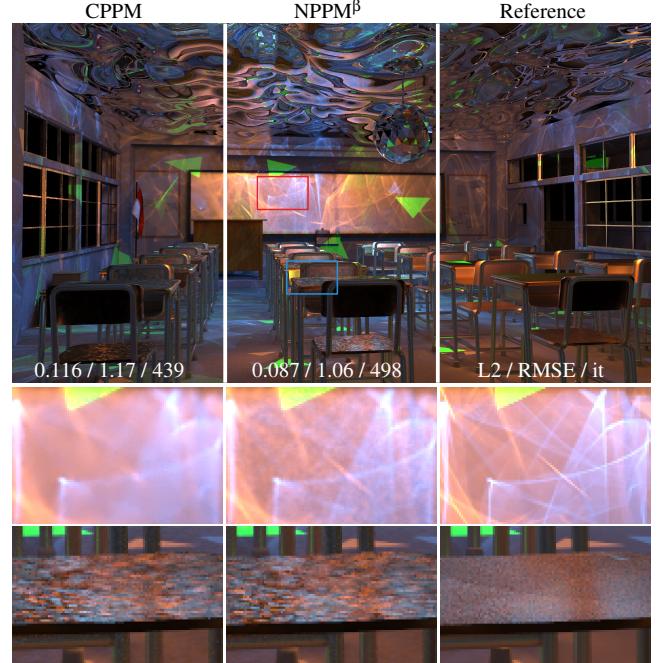


Figure 5: Equal-time comparison over five minutes of CPPM [LLZ*20] and NPPM $^\beta$ in the CLASSROOM scene. The scene contains several glossy materials, including the whiteboard, chairs, and desks. Glossy surfaces are challenging for photon mapping methods, particularly at grazing angles such as those observed on the chair seats and desks.

509 photons. Due to a more expressive kernel, our method can achieve
 510 improved convergence with significantly lower bias. We also in-
 511 clude results for our method NPPM $^\beta$. This more aggressive reduc-
 512 tion strategy reduces bias by employing smaller bandwidths. Un-
 513 like CPPM, this strategy does not rely on a statistical test; instead,
 514 the radius is reduced geometrically once a minimum photon count
 515 is reached. The rendered images in Fig. 6 show results compara-
 516 ble to CPPM. In scenes with depth of field effects, such as GLASS
 517 DOF, CPPM’s per-pixel statistical assumption breaks down, pro-
 518 ducing blurrier caustics compared to our method. However, in uni-
 519 formly lit regions, our method performs slightly worse than CPPM.

520 **Support for glossy materials.** We support a microfacet conductor
 521 BSDF (GGX) in our GPU implementation. Figure 5 demonstrates
 522 our method on glossy surfaces in the CLASSROOM scene. While
 523 our method exhibits some low-frequency noise on glossy surfaces,
 524 it outperforms CPPM, which tends to overblur sharp caustics. Graz-
 525 ing view angles produce noisy results across all our tests. Including
 526 roughness information in the DCV or combining our method with
 527 vertex connection and merging (VCM) [GKDS12] may be promis-
 528 ing directions for further improving results on such surfaces.

529 **Timings.** Table 1 quantifies the inference time of our NPPM ap-
 530 proach compared to SPPM on the ARTWARE and CRAB scenes.
 531 ARTWARE exhibits highly heterogeneous photon density, with cer-
 532 tain regions containing very high photon counts. In contrast, CRAB

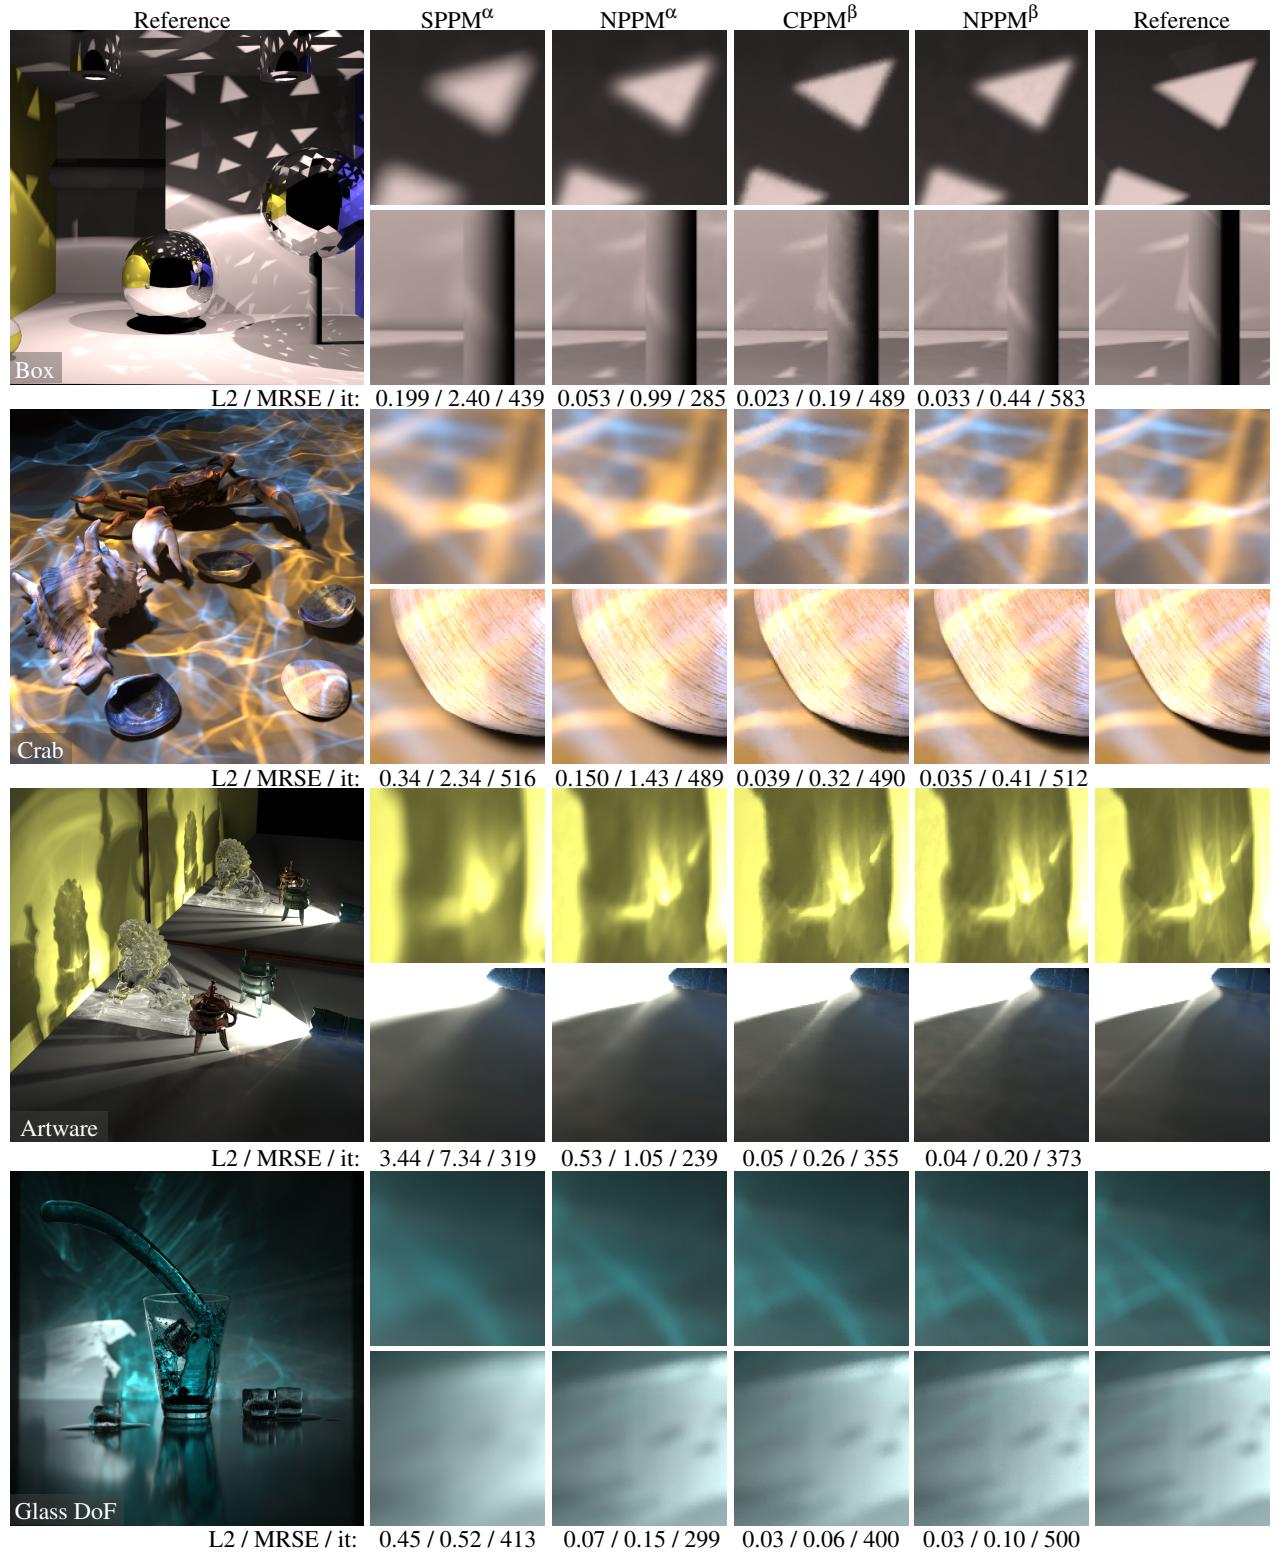


Figure 6: Comparison for NPPM with SPPM [HJ09] and CPPM [LLZ*20] at equal-time (30 seconds) with different radius reduction strategies. Our method produces sharper caustic details than SPPM. In some cases, CPPM tends to blur more complex caustic patterns (e.g., ARTWARE, GLASS DOF); however, its radius reduction strategy better preserves uniform, well-lit regions, which can lead to lower overall error.

Table 1: Performance breakdown between SPPM and our method (NPPM^α) for two representative scenes at different iteration counts (5 / 50 / 500 it.). All values are in milliseconds. Tracing time includes both eye and photon passes. Overhead is computed relative to the total SPPM time (Tracing + SPPM Eval.) compared to our method (Tracing + Enc. + Dec. + NPPM Eval.).

Scene	Tracing	SPPM Eval.	Enc. + Dec.	NPPM Eval.	Overhead (%)
ARTWARE	60.0	70.3 / 42.1 / 22.5	34.3 / 27.1 / 1.5	107.1 / 69.5 / 33.6	54.5% / 53.3% / 15.2%
CRAB	53.2	7.9 / 4.9 / 3.7	12.8 / 11.5 / 1.5	21.9 / 17.6 / 5.9	44.3% / 46.0% / 6.5%

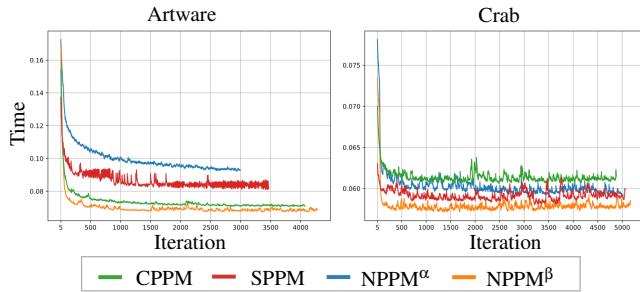


Figure 7: Per-iteration runtime for all methods on ARTWARE and CRAB. Runtimes are smoothed with an exponential moving average, and timings before the 5th iteration are excluded for improved readability.

533 features a more homogeneous and sparse photon distribution across 534 the scene. As a result, CRAB yields shorter range query and eval- 535 uation times than ARTWARE, as it produces a more uniform and 536 smaller workload on the GPU. During the first phase, in which we 537 encode 10% of the gather points and update the grid, our method 538 represents a 45–55% overhead. This phase lasts for only 50 itera- 539 tions, after which we simply read the DCV from the grid and de- 540 code it. This results in a final overhead ranging from 6–15%. Fig- 541 ure 7 provides a broader view of per-iteration cost during pro- 542 gressive rendering. For both NPPM^β and CPPM, the range query radius 543 decreases more rapidly, leading to faster evaluation and further re- 544 ducing overhead. However, the highly heterogeneous query sizes 545 introduced by CPPM’s radius reduction scheme can cause GPU 546 performance slowdowns. In contrast, our method makes more ef- 547 ficient use of GPU memory, enabling higher iteration counts under 548 the same time budget. Further optimizations are possible; for in- 549 stance, using a hash grid similar to Binder et al. [BFK20] or other 550 spatially accelerated data structures would likely reduce our over- 551 head further.

552 **Unbiased estimators** Explicit control over the kernel bandwidth 553 is required for unbiased density estimation via telescopic se- 554 ries [MBGJ22], which is naturally supported by our adaptive kernel 555 approach. For this application, we adapt our parametric kernel with

$$Q_\tau^{\text{Cone}}(\mathbf{d}) = \mathbf{1}_B(\mathbf{d}) \cdot (1 - \mathbf{d})^\tau / Z_{\text{cone}}, \quad (13)$$

556 so that it mimics a parametric cone kernel, yielding smaller asymp- 557 totic variance compared to constant kernels [MBGJ22]. We obtain 558 the normalization term $Z_{\text{cone}} = \frac{2\pi}{(\tau+1)(\tau+2)}$ by analytically integrat- 559 ing the kernel function. Our neural architecture is optimized to pro- 560 duce a suitable $\tau > 0.1$ for adaptive density estimation, in the same

563 submitted to EUROGRAPHICS 2026.

561 way as Q_Σ^{2DG} . This experiment also illustrates how our framework 562 can be extended to other parametric kernels, highlighting its poten- 563 tial for generalization.

564 **Figure 8** compares progressive estimators obtained with our 565 method and with standard density estimation, along with their un- 566 biased variants. We use the same random numbers to correlate the 567 stochastic radius selections and the photon map generations, since 568 unbiased variants often exhibit instability in convergence curves 569 due to rare unlikely very small radii sampled. The comparison 570 demonstrates that our method benefits both progressive and unbi- 571 ased estimators by reducing overall error. While unbiased estima- 572 tors show larger error than their biased counterparts, the resulting 573 images remain compatible with existing denoising methods. It is 574 important to note that, in theory, correlations between kernel es- 575 timations and local statistics during the grid construction phase may 576 introduce bias. However, we did not observe clear bias both visu- 577 ally and quantitatively.

Convergence analysis Figure 9 shows the evolution of L2 and MRSE for our methods, SPPM, and CPPM over rendering time. These results reinforce the trends observed in Figure 6: the introduction of depth of field reduces the performance gap between our approach and CPPM. Notably, our method consistently achieves lower error rates at shorter durations, which is advantageous for preview settings. At longer durations, our method delivers rendering quality that is on par with, or slightly below, that of CPPM on average.

5.1. Ablation study

Range-query versus K-NN Figure 10 shows a comparison between photon mapping, DPM [ZXJ*20] and our method with range-query and K -NN. We can see that adopting a progressive approach is important in order to reduce bias. In this scenario, DPM diverges since nothing in its design enforces the normalization of the kernel; consequently its consistency suffers. Our approach strongly reduces bias on sharp caustics while preserving smooth lighting in uniform density regions. Thus making it effective in range queries and progressive settings thanks to our normalized kernels.

Number of iterations for DCV averaging. While our method does not require online training during inference, DCV formation remains a major bottleneck since its cost grows linearly with the number of gathered photons. To mitigate this, our approach relies on spatial latent averaging (Section 4.4) and gatherpoint subsampling. Table 2 illustrates the diminishing impact of the number of

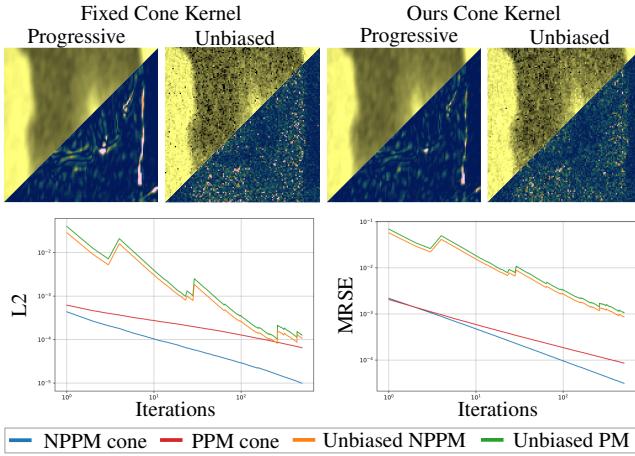


Figure 8: Comparison between our method using an adaptive cone kernel, SPPM [HJ09], and Unbiased PM [MBGJ22], both using fixed cone kernels. Our method reduces bias in sharp caustic regions and decreases fireflies in the unbiased case. Error maps show MRSE for progressive methods and SMAPE for unbiased counterparts.

averaging iterations on both error reduction and rendering time. Based on these results, we select 50 iterations as a practical compromise between computational overhead and overall accuracy.

Grid resolution. The resolution of the grid used to store the DCV spatially is an important hyperparameter. Low resolutions tend to over-blur the latent codes, resulting in poorly fitted kernels and visible bias in the final image. Conversely, higher resolutions increase memory consumption and slightly raise the cost of retrieval and updates. Table 3 summarizes the impact of grid resolution on image quality, showing diminishing returns beyond a certain resolution. In all our experiments, we use a 256^3 grid, which fits comfortably within the GPU memory. More optimized data structures, such as hash grids or tree-based layouts, could further improve performance while maintaining a comparable memory footprint.

The resolution of the grid used to store the DCV spatially is an important hyperparameter. Low resolutions tend to over-blur the latent codes, resulting in poorly fitted kernels and visible bias in the final image. Conversely, higher resolutions increase memory consumption and slightly raise the cost of retrieval and updates. Table 3 summarizes the impact of grid resolution on image quality, showing diminishing returns beyond a certain resolution. In all our experiments, we use a 256^3 grid, which fits comfortably within the GPU memory.

Figure 5 and Figure 11 demonstrate that the uniform grid can also handle larger scenes. For even larger environments, we recommend using adaptive data structures, such as sparse grids, octrees, or kD-trees constructed from gather-point statistics during a pre-processing stage. These structures would significantly improve the memory efficiency of DCV storage, although computation speed may be reduced due to the logarithmic traversal complexity. Exploring such data structures and implementing the corresponding DCV interpolation are left for future work.

Table 2: Effect of number of iterations for the DCV averaging for 100 passes in the ARTWARE with NPPM^β.

Iterations	MRSE ($\times 10^{-4}$)	MSE ($\times 10^{-5}$)	Time (s)
1 iter.	8.83	2.86	12.41
10 iter.	2.84	1.40	12.59
100 iter.	2.60	1.11	13.57

Table 3: Impact of grid resolution on rendering error for the ARTWARE scene using NPPM^α with 100 iterations. The 400^3 configuration corresponds to the maximum resolution that fits within our hardware.

Grid Resolution	MRSE ($\times 10^{-3}$)	MSE ($\times 10^{-4}$)
64^3	3.00	1.74
128^3	2.35	1.36
256^3	2.31	1.33
400^3	1.99	1.21

6. Conclusions

We have presented neural progressive photon mapping (NPPM), a learning-based framework that addresses kernel parameter selection challenges in density estimation through spatially adaptive, analytical kernels. Our lightweight neural model efficiently adapts to local photon distributions while maintaining compatibility with progressive rendering workflows and consistent estimators. NPPM represents the first learning-based progressive photon mapping method, providing a practical solution for high-quality caustic rendering that combines neural adaptivity with the controllability required for robust light transport simulation.

Limitations and future work. Although independent truncated Gaussians produce consistent estimators, our kernels actually depend on the photon distribution and are therefore correlated. As such, we cannot claim theoretical consistency for our estimator. In practice, this does not cause issue with NPPM^α, as shown in the convergence plots in Fig. 9.

For future work, we posit that incorporating material information into our photon features before encoding could potentially produce better-adjusted kernels. Extending our method to support more expressive parametric kernels, such as multi-modal or piecewise distributions, could also be particularly beneficial for volumetric photon rendering and guiding, where complex scattering patterns require sophisticated kernel representations to be well captured.

References

- [BFK20] BINDER N., FRICKE S., KELLER A.: Massively parallel path space filtering. In *International Conference on Monte Carlo and Quasi-Monte Carlo Methods in Scientific Computing* (2020), Springer, pp. 149–168. 5, 9
- [CKS*17] CHAITANYA C. R. A., KAPLANYAN A. S., SCHIED C., SALVI M., LEFOHN A., NOWROUZEZAHRAI D., AILA T.: Interactive reconstruction of Monte Carlo image sequences using a recurrent denoising autoencoder. 98:1–98:12. doi:10/gbxhcv. 2

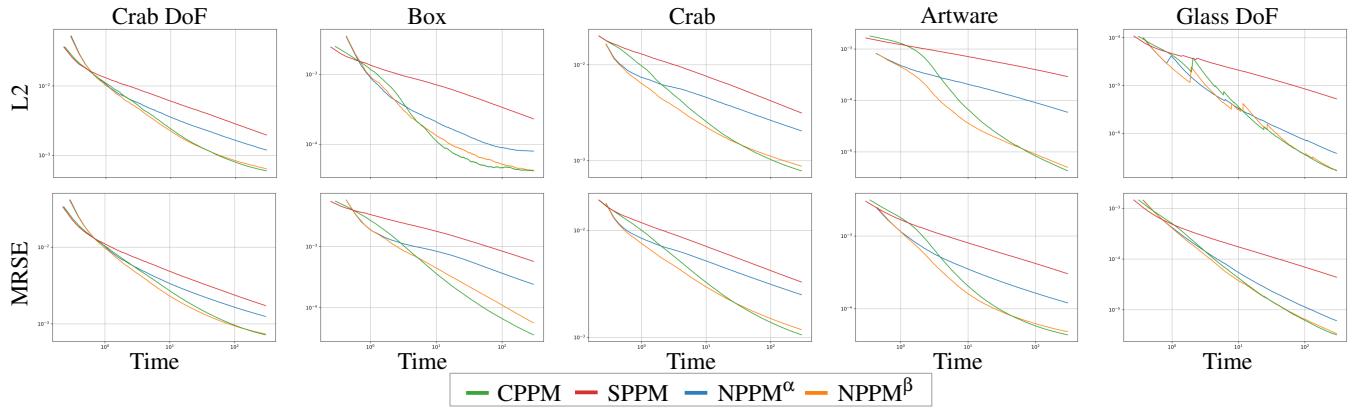


Figure 9: Convergence curves per time step (seconds) of our methods NPPM, SPPM [HJ09] and CPPM [LLZ*20] over five minutes. Our method with the APA radius reduction scheme NPPM $^\alpha$ exhibits consistent convergence and higher quality renders than SPPM. Although NPPM $^\beta$ demonstrates superior performance over CPPM in short-duration scenarios, the hypothesis testing scheme of CPPM leads to lower or equivalent long-term error rates.

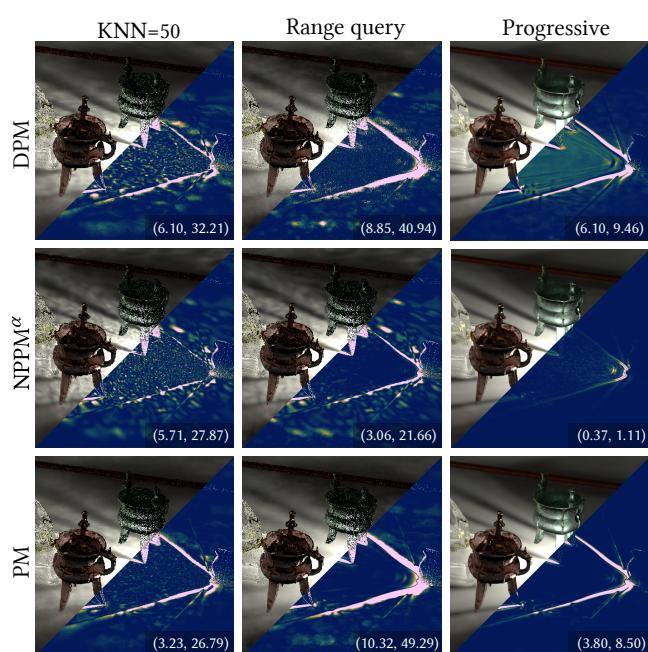


Figure 10: Comparison between our method, DPM [ZXJ*20] and SPPM [HJ09]. Neural methods outperform box kernel for single queries. However, due to its non-progressive design, DPM diverges when radius reduction is applied. We show the SMAPE in false color; numbers refer to L^2 and MRSE scores, respectively.

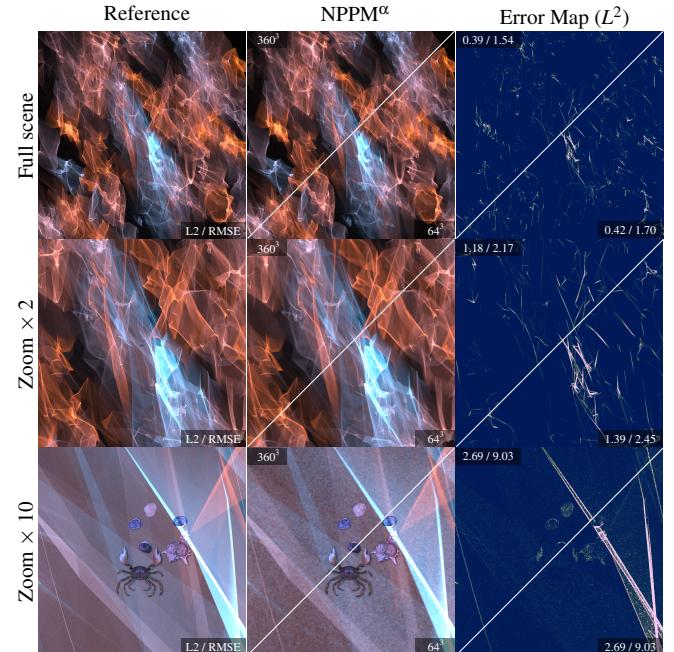


Figure 11: Equal-iteration comparison of different grid resolutions for a larger version of the CRAB scene across multiple zoom levels. The renderings were obtained with 2M photons per iteration and identical initial radii.

- 669 [CM21] CHOI H., MOON B.: Consistent post-reconstruction for progressive photon mapping. *Computer Graphics Forum* 40, 7 (2021), 121–130.
 doi:[10.1111/cgf.14406](https://doi.org/10.1111/cgf.14406). 2
- 670 [CWY11] CHEN J., WANG B., YONG J.-H.: Improved stochastic pro-
 671 gressive photon mapping with metropolis sampling. In *Computer Graph-
 672 ics Forum* (2011), vol. 30, Wiley Online Library, pp. 1205–1213. doi:
 673 [10.1111/j.1467-8659.2011.01979.x](https://doi.org/10.1111/j.1467-8659.2011.01979.x). 2
- 674 [DHC*21] DENG X., HAŠAN M., CARR N., XU Z., MARSCHNER S.:
 675 676

677 Path graphs: iterative path space filtering. *ACM Transactions on Graph-
 678 ics (Proceedings of SIGGRAPH Asia)* 40, 6 (2021). doi:[10.1145/3478513.3480547](https://doi.org/10.1145/3478513.3480547). 2

679 [DWL23] DONG H., WANG G., LI S.: Neural parametric mixtures for
 680 path guiding. In *ACM SIGGRAPH Conference Papers* (2023), pp. 1–10.
 doi:[10.1145/3588432.3591533](https://doi.org/10.1145/3588432.3591533). 2, 6

681 [EK20] ESTEVEZ A. C., KULLA C.: Practical caustics rendering with
 682 adaptive photon guiding. In *ACM SIGGRAPH Talks* (New York, NY,
 683 USA, 2020), Association for Computing Machinery. doi:[10.1145/3392734.3392740](https://doi.org/10.1145/3392734.3392740). 684 685

- 686 3388767.3407370. 2
 687 [FGW*24] FAN Z., GUO J., WANG Y., XIAO T., ZHANG H., ZHOU C.,
 688 CHEN Z., HONG P., GUO Y., YAN L.-Q.: Specular polynomials. *ACM
 689 Transactions on Graphics (Proceedings of SIGGRAPH)* 43, 4 (2024).
 690 doi:10.1145/3658132. 1
 691 [FWW*25] FAN Z., WANG C., WANG Y., LI B., GUO Y., YAN L.-Q.,
 692 GUO Y., GUO J.: Bernstein bounds for caustics. *ACM Transactions
 693 on Graphics (Proceedings of SIGGRAPH)* (2025). doi:10.1145/
 694 3731145. 1
 695 [GKDS12] GEORGIEV I., KŘIVÁNEK J., DAVIDOVIČ T., SLUSALLEK
 696 P.: Light transport simulation with vertex connection and merging.
 697 *ACM Transactions on Graphics (Proceedings of SIGGRAPH Asia)* 31,
 698 6 (2012). doi:10.1145/2366145.2366211. 2, 7
 699 [GPGSK18] GRITTMANN P., PÉRARD-GAYOT A., SLUSALLEK P.,
 700 KŘIVÁNEK J.: Efficient caustic rendering with lightweight photon map-
 701 ping. In *Computer graphics forum* (2018), vol. 37, Wiley Online Library.
 702 doi:10.1111/cgf.13481. 2
 703 [GRŠ*16] GRUSON A., RIBARDIÈRE M., ŠIK M., VORBA J., COZOT
 704 R., BOUATOUCH K., KŘIVÁNEK J.: A spatial target function for
 705 metropolis photon tracing. *ACM Transactions on Graphics (Proceed-
 706 ings of SIGGRAPH)* 36, 4 (2016). doi:10.1145/2963097. 2
 707 [HIT*24] HUANG J., IIZUKA A., TANAKA H., KOMURA T., KITA-
 708 MURA Y.: Online neural path guiding with normalized anisotropic spher-
 709 ical gaussians. *ACM Transactions on Graphics* 43, 3 (2024). doi:
 710 10.1145/3649310. 2, 6
 711 [HJ09] HACHISUKA T., JENSEN H. W.: Stochastic progressive pho-
 712 ton mapping. In *ACM Transactions on Graphics (Proceedings of SIG-
 713 GRAPH Asia)*. ACM New York, NY, USA, 2009. doi:10.1145/
 714 1618452.1618487. 1, 2, 3, 7, 8, 10, 11
 715 [HJ11] HACHISUKA T., JENSEN H. W.: Robust adaptive photon trac-
 716 ing using photon path visibility. *ACM Transactions on Graphics* 30, 5
 717 (2011). doi:10.1145/2019627.2019633. 2
 718 [HOJ08] HACHISUKA T., OGAKI S., JENSEN H. W.: Progressive photon
 719 mapping. *ACM Transactions on Graphics (Proceedings of SIGGRAPH
 720 Asia)* 27, 5 (2008). doi:10.1145/1409060.1409083. 2, 3, 6
 721 [HPJ12] HACHISUKA T., PANTALEONI J., JENSEN H. W.: A path space
 722 extension for robust light transport simulation. *ACM Transactions on
 723 Graphics (Proceedings of SIGGRAPH Asia)* 31, 6 (2012). doi:10.
 724 1145/2366145.2366210. 2
 725 [HY21] HUO Y., YOON S.-E.: A survey on deep learning-based monte
 726 carlo denoising. *Computational Visual Media* 7 (2021). doi:10.
 727 1007/s41095-021-0209-9. 2
 728 [IMF*21] İŞIK M., MULLIA K., FISHER M., EISENMANN J., GHARBI
 729 M.: Interactive monte carlo denoising using affinity of neural fea-
 730 tures. *ACM Transactions on Graphics (Proceedings of SIGGRAPH)* 40,
 731 4 (2021). doi:10.1145/3450626.3459793. 2
 732 [JC98] JENSEN H. W., CHRISTENSEN P. H.: Efficient simulation of light
 733 transport in scenes with participating media using photon maps. In *Pro-
 734 ceedings of the annual conference on Computer graphics and interactive
 735 techniques* (1998). doi:10.1145/280814.280925. 2
 736 [Jen95] JENSEN H. W.: Importance driven path tracing using the
 737 photon map. In *Rendering Techniques: Proceedings of the Euro-
 738 graphics Workshop* (1995), Springer, pp. 326–335. doi:10.1007/
 739 978-3-7091-9430-0_31. 2
 740 [Jen96] JENSEN H. W.: Global illumination using photon maps. In
 741 *Rendering Techniques* (Vienna, 1996), Springer, pp. 21–30. doi:10.
 742 1007/978-3-7091-7484-5_3. 1, 2, 3
 743 [JMLH01] JENSEN H. W., MARSCHNER S. R., LEVOY M., HANRA-
 744 HAN P.: A practical model for subsurface light transport. In *Proceedings
 745 of the Annual Conference on Computer Graphics and Interactive Tech-
 746 niques* (New York, NY, USA, 2001), Association for Computing Ma-
 747 chinery, p. 511–518. doi:10.1145/383259.383319. 2
 748 [JSR*22] JAKOB W., SPEIERER S., ROUSSEL N., NIMIER-DAVID
 749 M., VICINI D., ZELTNER T., NICOLET B., CRESPO M., LEROY
 750 V., ZHANG Z.: Mitsuba 3 renderer, 2022. URL: <https://mitsuba-renderer.org/>. 7
 751
 752 [JZJ08] JAROSZ W., ZWICKER M., JENSEN H. W.: The beam radi-
 753 ance estimate for volumetric photon mapping. In *ACM Transactions
 754 on Graphics (Proceedings of SIGGRAPH)* (2008). doi:10.1145/
 755 1401132.1401137. 2
 756 [KB15] KINGMA D. P., BA J.: Adam: A method for stochastic optimiza-
 757 tion. In *International Conference on Learning Representations (ICLR)*
 758 (2015), Bengio Y., LeCun Y., (Eds.). 6
 759 [KBG23] KERN R., BRÜLL F., GROSCH T.: Accelerating photon map-
 760 ping for hardware-based ray tracing. *Journal of Computer Graph-
 761 ics Techniques (JCGT)* 12, 1 (2023). URL: <http://jcgtr.org/published/0012/01/01/>. 3
 762
 763 [KD13] KAPLANYAN A. S., DACHSBACHER C.: Adaptive progressive
 764 photon mapping. *ACM Transactions on Graphics (Proceedings of SIG-
 765 GRAPH)* 32, 2 (2013). doi:10.1145/2451236.2451242. 2, 3
 766
 767 [KDB16] KELLER A., DAHM K., BINDER N.: Path space filtering. In
 768 *Monte Carlo and Quasi-Monte Carlo Methods*. Springer, 2016, pp. 423–
 436. doi:10.1007/978-3-319-33507-0_21. 2
 769
 770 [KGH*14] KŘIVÁNEK J., GEORGIEV I., HACHISUKA T., VĚVODA P.,
 771 ŠIK M., NOWROUZEZHRAI D., JAROSZ W.: Unifying points, beams,
 772 and paths in volumetric light transport simulation. *ACM Transactions
 773 on Graphics (Proceedings of SIGGRAPH)* 33, 4 (2014). doi:10.1145/
 774 2601097.2601219. 2
 775
 776 [KKLD23] KERBL B., KOPANAS G., LEIMKÜHLER T., DRETTAKIS G.:
 777 3d gaussian splatting for real-time radiance field rendering. *ACM Trans-
 778 actions on Graphics (Proceedings of SIGGRAPH)* 42, 4 (July 2023).
 779 doi:10.1145/3592433. 4
 780
 781 [KPB21] KOBYZEV I., PRINCE S. J., BRUBAKER M. A.: Normalizing
 782 flows: An introduction and review of current methods. *IEEE transac-
 783 tions on pattern analysis and machine intelligence* 43, 11 (2021), 3964–3979.
 784 doi:10.1109/TPAMI.2020.2992934. 2
 785
 786 [KW14] KINGMA D. P., WELLING M.: Auto-encoding variational
 787 bayes. *International Conference on Learning Representations (ICLR)*
 788 (2014). 6
 789
 790 [KWX*16] KANG C.-M., WANG L., XU Y.-N., MENG X.-X., SONG
 791 Y.-J.: Adaptive photon mapping based on gradient. *Journal of Com-
 792 puter Science and Technology* 31 (2016), 217–224. doi:10.1007/
 793 s11390-016-1622-x. 2
 794
 795 [KZ11] KNAUS C., ZWICKER M.: Progressive photon mapping: A prob-
 796 abilistic approach. *ACM Trans. Graph.* 30, 3 (2011). doi:10.1145/
 797 1966394.1966404. 2, 3
 798
 799 [LHL*24] LITALIEN J., HAŠAN M., LUAN F., MULLIA K., GEORGIEV
 800 I.: Neural product importance sampling via warp composition. In
 801 *ACM SIGGRAPH Asia Conference Papers* (2024). doi:10.1145/
 802 3680528.3687566. 2
 803
 804 [LLZ*20] LIN Z., LI S., ZENG X., ZHANG C., JIA J., WANG G.,
 805 MANOCHA D.: Cppm: chi-squared progressive photon mapping. *ACM
 806 Transactions on Graphics (Proceedings of SIGGRAPH Asia)* 39, 6 (Nov.
 807 2020). doi:10.1145/3414685.3417822. 1, 2, 3, 7, 8, 11
 808
 809 [MŽ1] MÜLLER T.: tiny-cuda-nn, 2021. URL: <https://github.com/NVlabs/tiny-cuda-nn>. 7
 810
 811 [MBGJ22] MISSO Z., BITTERLI B., GEORGIEV I., JAROSZ W.: Un-
 812 biassed and consistent rendering using biased estimators. *ACM Trans-
 813 actions on Graphics (Proceedings of SIGGRAPH)* 41, 4 (2022). doi:
 814 10.1145/3528223.3530160. 2, 4, 9, 10
 815
 816 [MMR*19] MÜLLER T., McWILLIAMS B., ROUSSELLE F., GROSS
 817 M., NOVÁK J.: Neural importance sampling. *ACM Transactions on
 818 Graphics (Proceedings of SIGGRAPH)* 38, 5 (2019). doi:10.1145/
 819 3341156. 2

- 810 [PAMM19] PERROT R., AVENEAU L., MORA F., MENEVEAUX D.:
 811 Photon mapping with visible kernel domains. *The Visual Computer* 35
 812 (2019), 707–720. doi:[10.1007/s00371-018-1505-y](https://doi.org/10.1007/s00371-018-1505-y). 2
- 813 [PGM*19] PASZKE A., GROSS S., MASSA F., LERER A., BRADBURY
 814 J., CHANAN G., KILLEEN T., LIN Z., GIMELSHEIN N., ANTIGA L.,
 815 DESMAISON A., KOPF A., YANG E., DEVITO Z., RAISON M., TE-
 816 JANI A., CHILAMKURTHY S., STEINER B., FANG L., BAI J., CHIN-
 817 TALA S.: *PyTorch: An Imperative Style, High-Performance Deep Learn-
 818 ing Library.* 2019, pp. 8024–8035. doi:[10.5555/3454287.3455008](https://doi.org/10.5555/3454287.3455008). 7
- 820 [QSH*15] QIN H., SUN X., HOU Q., GUO B., ZHOU K.: Unbiased
 821 photon gathering for light transport simulation. *ACM Transactions on
 822 Graphics (Proceedings of SIGGRAPH Asia)* 34, 6 (2015). doi:[10.1145/2816795.281811](https://doi.org/10.1145/2816795.281811). 2, 4
- 823 [QSMG17] QI C. R., SU H., MO K., GUIBAS L. J.: Pointnet: Deep
 824 learning on point sets for 3d classification and segmentation, 2017.
 825 URL: <https://arxiv.org/abs/1612.00593>, arXiv:1612.
 826 00593. 2, 3
- 827 [Ram14] RAMA H.: Fast fixed-radius nearest neighbors: Interactive
 828 million-particle fluids. *GPU Technology Conference* (2014). 7
- 829 [SFES07] SCHJØTH L., FRISVAD J. R., ERLEBEN K., SPORRING J.:
 830 Photon differentials. In *Proceedings of the 5th international conference
 831 on Computer graphics and interactive techniques in Australia and South-
 832 east Asia* (2007), pp. 179–186. doi:[10.1145/1321261.1321293](https://doi.org/10.1145/1321261.1321293).
 833 2
- 834 [ŠK19] ŠIK M., KRIVÁNEK J.: Implementing one-click caustics in
 835 corona renderer. In *Proceedings of EGSR (Industry Track)* (2019),
 836 pp. 61–67. doi:[10.2312/sr.20191221](https://doi.org/10.2312/sr.20191221). 2
- 837 [Vea98] VEACH E.: *Robust Monte Carlo methods for light transport sim-
 838 ulation.* Stanford University, 1998. 2
- 839 [vOHK16] ŠIK M., OTSU H., HACHISUKA T., KRIVÁNEK J.: Ro-
 840 bust light transport simulation via metropolised bidirectional estimators.
 841 *ACM Transactions on Graphics (Proceedings of SIGGRAPH Asia)* 35, 6
 842 (2016). doi:[10.1145/2980179.2982411](https://doi.org/10.1145/2980179.2982411). 2
- 843 [WGGH20] WEST R., GEORGIEV I., GRUSON A., HACHISUKA T.:
 844 Continuous multiple importance sampling. *ACM Transactions on
 845 Graphics (Proceedings of SIGGRAPH)* 39, 4 (2020). doi:[10.1145/3386569.3392436](https://doi.org/10.1145/3386569.3392436). 2
- 846 [WGH22] WEST R., GEORGIEV I., HACHISUKA T.: Marginal multi-
 847 ple importance sampling. In *ACM SIGGRAPH Asia Conference Papers*
 848 (2022). doi:[10.1145/3550469.3555388](https://doi.org/10.1145/3550469.3555388). 2
- 849 [ZGJ20] ZELTNER T., GEORGIEV I., JAKOB W.: Specular manifold
 850 sampling for rendering high-frequency caustics and glints. *ACM Trans-
 851 actions on Graphics (Proceedings of SIGGRAPH)* 39, 4 (2020), 149–1.
 852 doi:[10.1145/3386569.3392408](https://doi.org/10.1145/3386569.3392408). 1
- 853 [ZXJ*20] ZHU S., XU Z., JENSEN H. W., SU H., RAMAMOORTHI R.:
 854 Deep kernel density estimation for photon mapping. *Computer Graphics
 855 Forum* 39, 4 (2020), 35–45. doi:[10.1111/cgf.14052](https://doi.org/10.1111/cgf.14052). 2, 3, 4, 5,
 856 7, 9, 11
- 857 [ZXS*21a] ZHU S., XU Z., SUN T., KUZNETSOV A., MEYER M.,
 858 JENSEN H. W., SU H., RAMAMOORTHI R.: Hierarchical neural recon-
 859 struction for path guiding using hybrid path and photon samples. *ACM
 860 Transactions on Graphics (Proceedings of SIGGRAPH)* 40, 4 (2021).
 861 doi:[10.1145/3450626.3459810](https://doi.org/10.1145/3450626.3459810). 2
- 862 [ZXS*21b] ZHU S., XU Z., SUN T., KUZNETSOV A., MEYER M.,
 863 JENSEN H. W., SU H., RAMAMOORTHI R.: Photon-driven neural re-
 864 construction for path guiding. *ACM Transactions on Graphics* 41, 1
 865 (2021). doi:[10.1145/3476828](https://doi.org/10.1145/3476828). 2